

Percipient - course Podssummary

1. Identify what a Pod is & Explain its crucial role in Kubernetes

What is a Pod?

A Pod is the smallest deployable and schedulable unit in Kubernetes.

It acts as a wrapper around:

- One container (most common case), or
- Multiple tightly coupled containers that must run together

Key characteristics

- **Shared execution environment** → containers inside a pod share:
 - Storage volumes
 - Network namespace (same IP & port space)
 - IPC (inter-process communication)
- Pods represent a **logical host** for an application.
- Containers inside a pod are always:
 - **Co-located** (run on the same node)
 - Co-scheduled
 - **Share resources** and are often complementary (e.g., Web server + sidecar)

Why Pods are crucial

- They define how an application runs in Kubernetes.
- They act as the **basic building block** upon which higher-level objects (Deployments, ReplicaSets, DaemonSets, Jobs) operate.
- They group containers that belong together and need to communicate efficiently.
- Each pod gets its **own unique IP address**, simplifying networking.
- Support init containers and ephemeral containers for:
 - Initialization sequences
 - Debugging

How Pods behave

- Pods are **ephemeral**: they don't self-heal.
- If a pod fails, Kubernetes **creates a new pod** rather than fixing the old one.
- Controlled by workload resources (Deployments, Jobs, etc.) using **PodTemplates**.

✓ 2. Identify the phases of a Pod's life & Describe its lifecycle

Pod Lifecycle Phases

Kubernetes tracks pod progress through the following **Pod phases**:

1. Pending
 - Pod is accepted by the cluster but containers haven't started yet.
 - Includes time for scheduling + image downloads.
2. Running
 - Pod has been scheduled to a node.
 - All containers created and at least one is running.
3. Succeeded
 - All containers terminated successfully and will not restart.
4. Failed
 - All containers terminated, and **at least one container failed**.
5. Unknown
 - API server can't communicate with the kubelet on that node.

Container States inside a Pod

Each container has its own state:

1. **Waiting** → preparing to start (pulling image, setup)
2. **Running** → container is operating normally
3. **Terminated** → completed or failed (includes exit code & reason)

Restart Policies

Defined in pod spec → affects what Kubernetes does if containers exit:

- **Always** → restart regardless of success/failure
- **OnFailure** → restart only if container exits with failure
- **Never** → never restart the container

(Example: batch jobs → Never/OnFailure, web servers → Always)

Pod Conditions

PodStatus contains conditions that indicate readiness:

- **PodScheduled** → pod has been assigned to a node
- **Initialized** → all init containers completed
- **ContainersReady** → all containers are ready
- **Ready** → pod is ready to serve traffic

Each condition has:

- type, status (True/False/Unknown), lastProbeTime, lastTransitionTime, reason, message

3. Identify how resources can be allocated & restricted per-container

Kubernetes allows you to control **CPU and memory** for each container using:

Resource Requests

- The **minimum guaranteed** amount of CPU/memory the container needs.
- Scheduler uses it to place pods on nodes.
- Example:

resources:

- requests:
- memory: "50Mi"
-

Resource Limits

- The **maximum** amount of CPU/memory a container is allowed to use.
- Prevents resource hogging.
- If a container exceeds:
 - Memory limit → OOMKill (Out Of Memory Kill)
 - CPU limit → throttled

Example:

resources:

limits:

memory: "250Mi"

Behaviour Example

- If container tries to allocate **more memory than the limit**, Kubernetes kills it:
 - Pod status becomes **CrashLoopBackOff**
 - Reason: **OOMKilled**

5. Kubernetes Namespaces – Summary

What is a Namespace?

A **namespace** is a logical partition inside a Kubernetes cluster that allows multiple *virtual clusters* within the same physical cluster.

Key Points

- They **scope** resource names (names must be unique inside a namespace).
- Resources cannot belong to more than one namespace.
- Namespaces cannot be nested.
- Best used when:
 - Many users
 - Spread across multiple teams/projects
(NOT needed for small teams.)

When *not* to use namespaces

- Don't create namespaces for different software versions → use **labels** instead.

Built-in Kubernetes Namespaces

1. **default** — for objects without a specified namespace
2. **kube-system** — for system/cluster components
3. **kube-public** — publicly readable by all (even unauthenticated users)
4. **kube-node-lease** — stores node heartbeat info for health checks

Commands

- List namespaces:
kubectl get namespace

Naming caution

- Don't start namespace names with **kube-** (reserved for system use)

✓ 6. Limit Ranges — Summary (Kubernetes Policies)

What is a Kubernetes Policy?

A set of rules enforced by the cluster to control how resources behave.

LimitRange (Policy Type)

A LimitRange limits:

- Minimum and maximum **memory/cpu** a container can use.
- Can also set **default requests/limits** when not defined in a pod.

Why needed?

By default, containers have **no resource constraints**, which risks cluster exhaustion.

How LimitRange works

Placed **inside a namespace** → applies to all pods in that namespace.

Kubernetes checks:

1. If pod does NOT specify memory requests/limits → apply defaults from LimitRange.
2. If pod specifies its own values → Kubernetes verifies:
 - request \geq min
 - limit \leq max
 - otherwise: Forbidden error

Example Behavior

- Pod without memory limits → gets default limit and request from LimitRange
- Pod with valid custom limits → allowed
- Pod exceeding max (e.g., 800Mi > 750Mi) → **creation denied**

7. ReplicaSets — Summary

What is a ReplicaSet?

A controller that ensures a stable number of identical pod replicas are running at all times.

Functions

- Self-healing (recreate pods if deleted)
- Scaling (increase/decrease replicas)
- Maintain desired number of pods

Key Concepts

- Part of the apps/v1 API group
- Used under the hood by **Deployments**
- Has:
 - **selector** → match labels to identify pods
 - **replicas** → number of desired pod copies
 - **template** → pod specification

ownerReferences

Each pod created by a ReplicaSet has an **ownerReference** linking it back to its ReplicaSet.

Acquiring pods

If a pod:

- has no ownerReference
- AND matches the ReplicaSet's selector
→ the ReplicaSet will **adopt** it.

Usage

We rarely manage ReplicaSets directly—Deployments typically handle them.

✓ 8. DaemonSets — Summary

What is a DaemonSet?

A controller that ensures **one pod per node** (or more, depending on selectors).

Why used?

To run system-level services on each node such as:

- Logging agents
- Monitoring agents
- Storage daemons

Behavior

- When a node is added → a pod is created on that node
- When a node is removed → the pod is deleted
- When DaemonSet is deleted → its pods are removed (unless cascade=false)

Accessing DaemonSet pods

Methods:

- **Push model** (pods push logs to a central service)
- NodeIP + HostPort
- DNS (headless service)
- **Regular service** (load balances randomly across nodes)

YAML requirements

- apiVersion: **apps/v1**
- kind: **DaemonSet**
- metadata: name, labels
- spec includes:
 - selector
 - template (pod spec)
- pod RestartPolicy must be **Always**

Scheduling

DaemonSet pods are **not created by the default scheduler**—they're created by the DaemonSet controller.

✓ 9. Kubernetes Deployments — Summary

What is a Deployment?

A higher-level controller used for declarative updates of pods and ReplicaSets.

Key Uses

- Rollout new ReplicaSets
- Roll back to previous versions
- Scale pods up/down
- Update pod template specs (new desired state)
- Clean up old ReplicaSets

Example use cases

- Rolling updates (zero-downtime)
- Versioned rollbacks
- Scaling to meet traffic demand

YAML Anatomy

- apiVersion: **apps/v1**
- kind: Deployment
- metadata: name
- spec:
 - replicas: number of desired pods
 - selector: matchLabels
 - template:
 - metadata: labels
 - spec: container definitions (image, ports, etc.)

Relationship

Deployment → manages ReplicaSet → manages Pods

✓ 11. Kubernetes Services — Key Concepts

Why Services?

- Pods are *ephemeral* → their IPs change whenever pods restart.
- You cannot rely on pod IPs.
- A **Service** gives a **stable virtual IP (ClusterIP)** that points to a logical set of pods.

◆ What is a Service?

A Service is:

- An **abstraction** that groups pods using **label selectors**.
- Exposes these pods as a **network service**.
- Ensures consistent connectivity even when pods die and recreate.

Example:

Frontend → Service → Backend pods
(Frontend never needs to know pod IPs.)

◆ How Kubernetes Service Discovery Works

1. DNS-based discovery (recommended)

- Internal DNS (like **CoreDNS**) watches the API server.
- Creates DNS records for each Service.
- Every Pod can access services by name:
`my-service.my-namespace.svc.cluster.local`

2. Environment variables (older method)

- Kubelet injects env vars for each service into pods.
- Works only if Service exists *before* Pods are created.

◆ Service YAML Basics

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: svcApp
  ports:
    - port: 80      # service port inside cluster
      targetPort: 9000 # container port
      protocol: TCP
```

◆ Services Without Selectors

Used when:

- Backend is **external DB** or external service.
- Service points to another namespace/cluster.
- Partial migration (some pods in Kubernetes, some outside).

Here YOU manually create **Endpoints** object.

◆ Headless Services (clusterIP: None)

Used when:

- You don't need load balancing
- You want DNS to return **pod IPs directly**
- Useful for:
 - StatefulSets
 - Databases (e.g., Cassandra, ZooKeeper)

◆ Service Types

Type	Description	Use Case
ClusterIP (default)	Internal-only access	Backend services
NodePort	Accessible on <code><NodeIP>:<NodePort></code>	Simple external access
LoadBalancer	Cloud provider LB	Public access in cloud
ExternalName	CNAME to external DNS	External service reference

✓ 12. Exploring Services in Action — Demo Summary

Goal:

Create a **NodePort Service** exposing an nginx pod.

Steps:

1. Created demo namespace:

```
kubectl create ns demo
```

- 2.

3. Created a Pod from YAML (`nginx:1.19.0`).

4. Exec into Pod:

```
kubectl exec nginx-pod -n demo -it -- /bin/sh
```

5.

- Custom index.html was added inside container.

6. Applied the Service YAML:

```
kind: Service
7. type: NodePort
8. ports:
9. - port: 80
10. targetPort: 80
11. nodePort: 31000
12. selector:
13. app: nginx
14.
```

15. Found Node IP:

```
kubectl get nodes -o wide
16.
```

17. Accessed service in browser:

```
http://<NodeIP>:31000
18.
```

Result: Custom Nginx page appears using NodePort access.

13. Identifying Ingress — Manage External Traffic

What is an Ingress?

- A Kubernetes object that routes external HTTP/HTTPS traffic into the cluster.
- Uses **Ingress rules** to route URLs → Services.
- Needs an **Ingress Controller** to function (e.g., Nginx Ingress Controller).

Ingress Requirements

1. Ingress Controller

In minikube:

```
minikube addons enable ingress
```

2.

3. Ingress Object

Created using YAML manifest.

Demo Architecture

Two pods + services:

Pod	Service	Message	URL Path
app-1	app1-svc	greetings from app1	/path1
app-2	app2-svc	greetings from app2	/path2

Ingress YAML Example

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: simple-ingress
```

```
spec:
```

```
rules:  
- host: domain.local  
  http:  
    paths:  
      - path: /path1  
        pathType: Prefix  
        backend:  
          service:  
            name: app1-svc  
            port:  
              number: 5678  
      - path: /path2  
        pathType: Prefix  
        backend:  
          service:  
            name: app2-svc  
            port:  
              number: 5678
```

◆ Checking Ingress

Watch for IP assignment:

```
kubectl get ing simple-ingress -o wide -w
```

Minikube ingress exposes:

```
minikube ip
```

Map domain locally (in Windows hosts file):

```
172.17.20.26 domain.local
```

◆ Access Ingress

In browser:

- `http://domain.local/path1 → app1 response`
- `http://domain.local/path2 → app2 response`

 Final Key Differences to Remember

Feature	Service	Ingress
Purpose	Internal service discovery & stable networking	HTTP routing from outside
Gateway?	No	Yes (reverse proxy)
Works on	L3/4 (ports)	L7 (HTTP/HTTPS)
Needs controller?	No	Yes
Supports path-based routing?	No	Yes